Norynthe.

AI TRUST INFRASTRUCTURE

Norynthe Investor Packet

Founder-built AI evaluation OS for smaller-model review today; dedicated lab and compute capacity unlock larger-model benchmarking, trust scoring, and pilot-ready evidence.

80
SAMPLE SCORE

Evidence behind the score.

The current OS already presents console workflow, assessment records, report surfaces, and capital logic; compute is the next constraint.

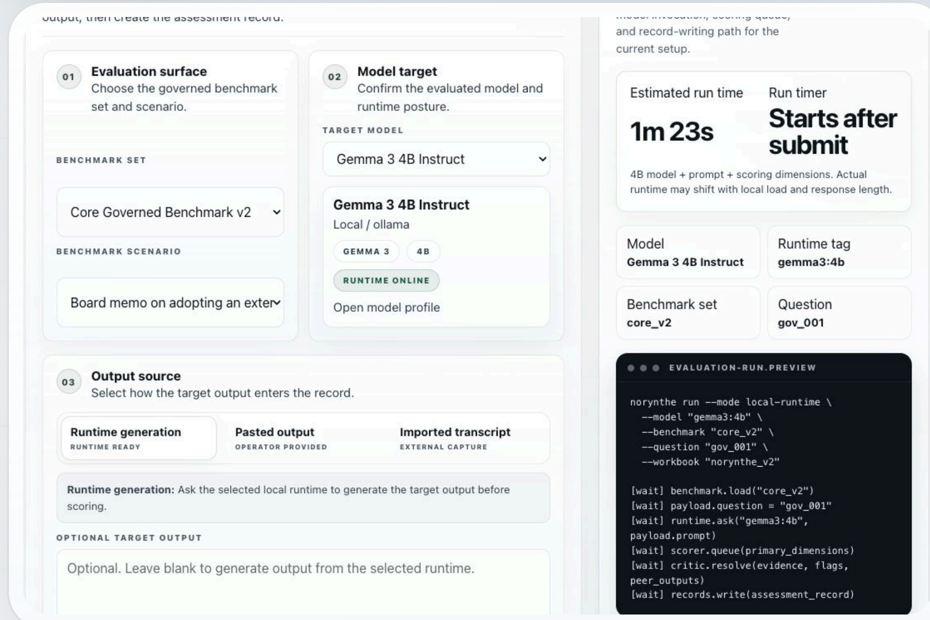

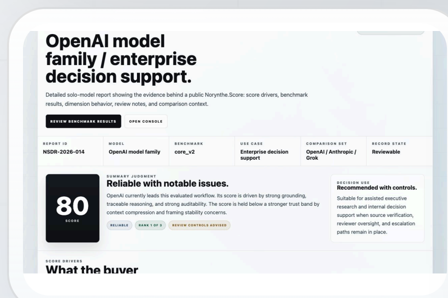

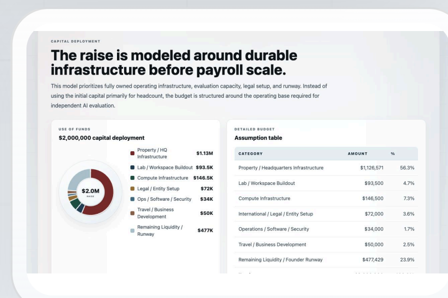

RAISE TARGET
$2.0M
Dedicated lab and compute capacity.

PRODUCT STATE
Evaluation OS live
Smaller-model evaluation, console, records, reports, and investor surfaces.

COMMERCIAL MODEL
Paid evidence
Reports, methodology, comparisons, and subscriptions.

01 / EXECUTIVE SNAPSHOT

Norynthe operates the external trust layer for AI systems.

Norynthe.Score is the public signal. The founder-built evaluation OS exists today for smaller-model review; dedicated compute unlocks larger-model standing, score detail, and institution-ready reporting.

External evidence becomes the buying standard.

Enterprises need a way to compare AI systems outside model-owner claims, vendor demos, and narrow benchmark headlines. Norynthe creates a governed evidence record that can travel through procurement, risk, and executive decision processes.

1. **Working evaluation OS**

   The current OS evaluates smaller models and produces console, assessment record, report, and investor evidence surfaces.

2. **Independent score layer**

   The score is built from governed benchmark assets, not model-owner marketing claims.

3. **Compute bottleneck unlock**

   The $2M round funds dedicated lab compute for larger-model evaluation, evidence hardening, and enterprise validation.

01 Benchmark assets

Governed prompts, rubrics, test dimensions, and evaluation conditions.

02 Norynthe Model

Evaluator, scoring logic, confidence, blocker criteria, and judgment rules.

03 Assessment records

Versioned model standing, score detail, review state, findings, and flags.

04 Enterprise reports

Buyer-readable evidence for procurement, governance, risk, and operations.

02 / COMPANY THESIS

Independent scoring becomes infrastructure.

Most AI evaluation collapses into demos, brand perception, leaderboard headlines, or provider-controlled tests. Norynthe starts from a different premise: the workflow exists, and the next bottleneck is compute capacity.

CATEGORY THESIS

Trust moves outside the model.

Norynthe is building the record system enterprises can use when model claims are not enough.

Benchmark bank	Scoring engine
Assessment record	Report surface

01 **AI choice is now an operating decision.**

Model selection is moving into workflows with procurement, policy, legal, risk, and operational consequences.

02 **Existing evaluation is too vendor-shaped.**

Leaderboards and demos are useful, but they do not create a neutral record buyers can reuse across departments.

03 **A governed evidence record becomes reusable infrastructure.**

The score creates comparison pressure; the record and report create the paid evidence layer behind it.

04 **The moat compounds through evidence density.**

Benchmark governance, scoring logic, reviewer state, and buyer-facing reports harden together as more evaluations run.

03 / PRODUCT PROOF

The evaluation OS exists today.

The investor preview shows the working console behind Norynthe.Score: smaller-model evaluation launch, runtime trace, reviewer queue, stored assessment records, benchmark governance, and report-ready evidence.

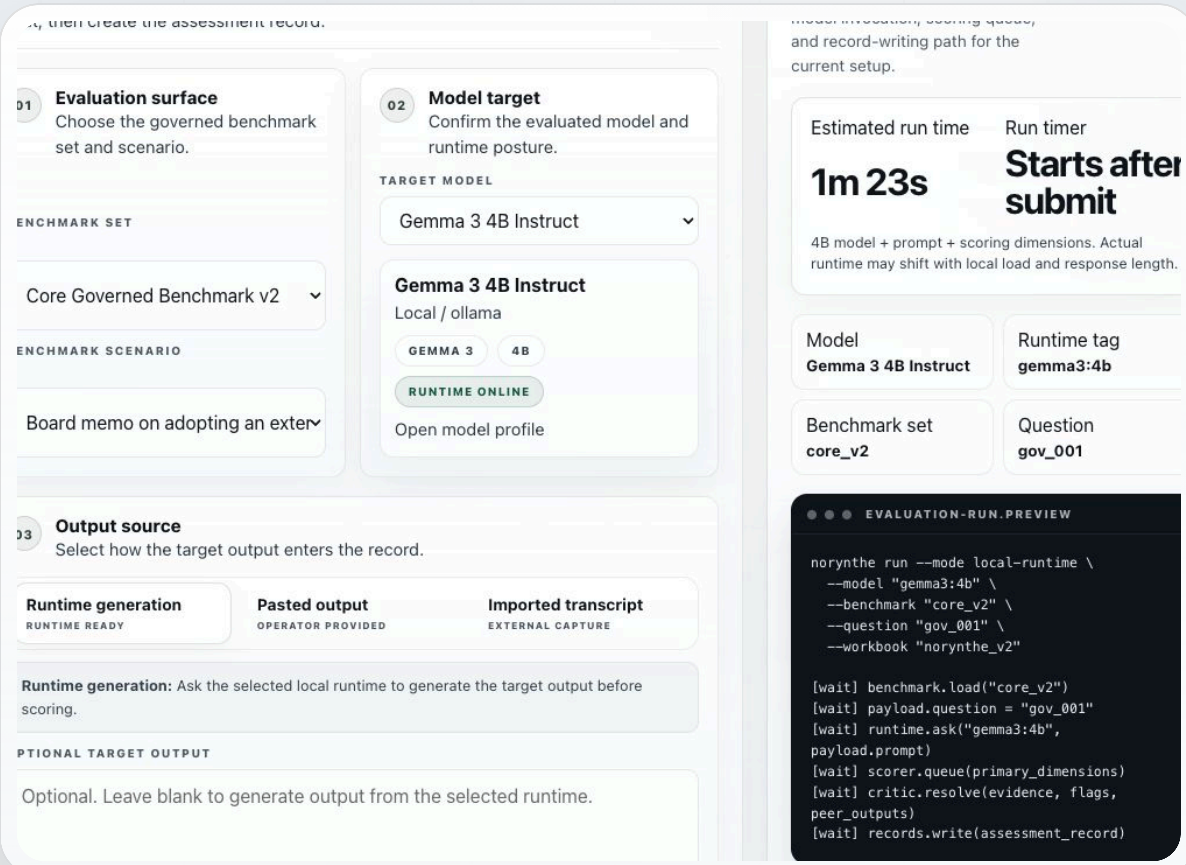

CURRENT PROOF

Working evaluation workflow before dedicated compute expansion.

Launch
Benchmark, target model, runtime posture, and evaluation conditions.

Score
Evaluator logic, dimensions, evidence checks, flags, and confidence.

Record
Assessment state, reviewer posture, persisted findings, and caveats.

Report
Buyer-facing output generated from the underlying evidence record.

04 / BUSINESS MODEL

Public trust scores create demand. Paid evidence access creates revenue.

Norynthe sells access to the evidence behind independent AI evaluation. The score is the entry point; revenue comes when teams need the methodology, comparisons, records, and reports behind it.

PUBLIC LAYER

Norynthe.Score

A clear public signal that creates comparison pressure and gives buyers a reason to ask for the evidence behind the ranking.

EVIDENCE LAYER

Reports and records

Paid access to full evaluation findings, methodology detail, score context, evidence trails, and decision-ready summaries.

RECURRING LAYER

Enterprise access

Subscriptions, custom evaluations, vendor comparisons, advisory review, API access, and licensing as trust demands repeat.

Full evaluation reports	Score methodology	Vendor comparisons
Evidence trails	Governance documentation	Custom evaluations

REVENUE LOGIC

The score creates the reason to ask. The paid product is the evidence behind the score: methodology, comparisons, documentation, and decision support.

05 / ENTERPRISE USE CASES

Norynthe turns model trust into a decision workflow.

The enterprise need is not simply more model access. It is better judgment about which systems should be used, under what constraints, with what evidence, and with what level of review.

COMPARE

Which model has standing?

Rank systems under a shared benchmark surface instead of vendor-specific claims.

EXPLAIN

Why did one system outrank another?

Expose dimensions, evidence, flags, and score context behind the public result.

ESCALATE

Where should human review remain mandatory?

Identify blocker criteria, caveats, risk posture, and sensitive-use constraints.

OPERATIONALIZE

How should the selected system be used?

Translate evaluation evidence into procurement, governance, and deployment language.

Model selection	Compare vendors under the same benchmark surface.
Procurement review	Reduce dependence on vendor-defined evidence.
Internal governance	Make trust judgments reviewable across teams.
Sensitive workflows	Apply blocker criteria and escalation logic.
Operational deployment	Decide where constraints and human oversight belong.

USE-CASE READ

The product is strongest where decisions need reviewable evidence, shared language, and a standard outside the model vendor.

06 / RAISE PLAN

$2M to remove the compute bottleneck.

The round moves the company from smaller-model proof into a controlled evaluation lab with the compute capacity required for larger-model trust scoring and enterprise validation.

TARGET RAISE

$2.0M

Infrastructure-first capital for an owned operating base, dedicated lab compute, legal setup, security, and founder-led validation.

CAPITAL READ

The plan prioritizes compute capacity, durable lab infrastructure, and validation before broad payroll expansion.

01 **Secure the base**

Property, legal setup, operating environment, and investor-ready governance foundation.

02 **Build the lab**

Workspace, controlled compute, storage, security, and repeatable evaluation infrastructure.

03 **Harden the model**

Benchmark governance, scoring logic, review-state discipline, and assessment records.

04 **Validate with buyers**

Enterprise workflows, pilots, paid report access, custom evaluation demand, and subscriptions.

07 / USE OF FUNDS

Capital goes into compute capacity before payroll scale.

The allocation prioritizes durable operating infrastructure, dedicated lab compute, larger-model evaluation capacity, legal setup, controlled operations, business development, and focused founder-led execution.

TOTAL RAISE

$2.0M

The use-of-funds view is intentionally literal: every bar is proportional to the full $2M raise, so the dominant property and operating-reserve allocations remain visually honest.

Owned base	56.3%
Execution reserve	23.9%
Lab and compute	12.0%

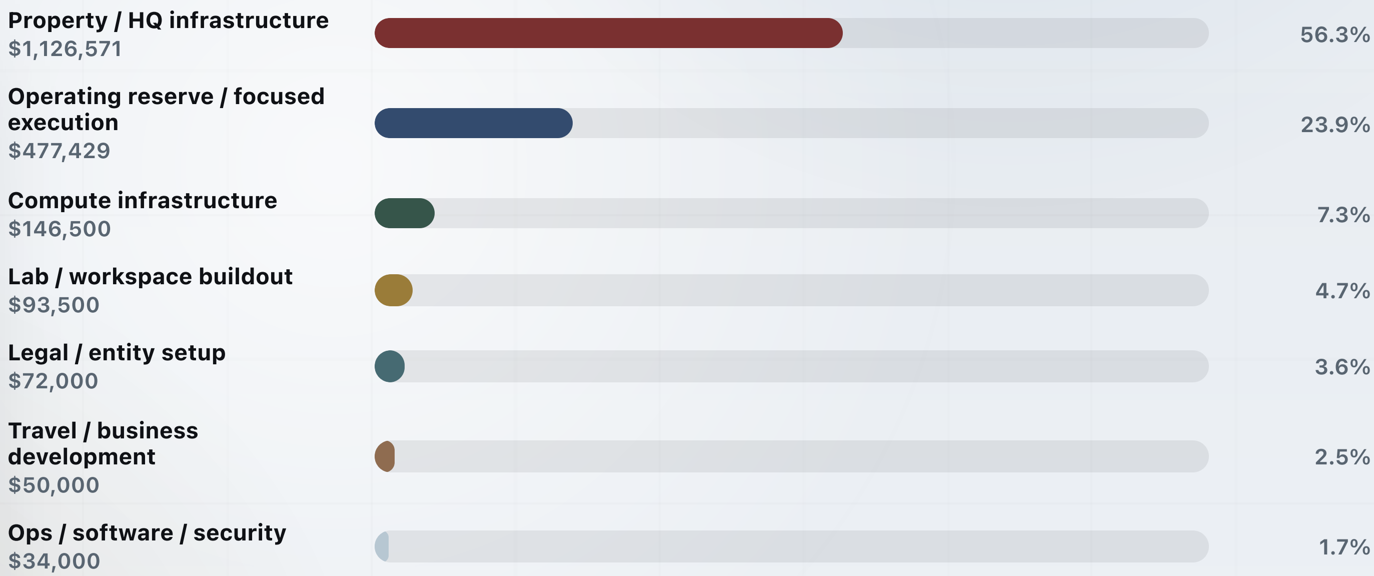

08 / FINANCIAL MODEL

A 24-month base scenario after a six-month ramp.

The forecast models enterprise subscriptions, custom evaluations, reports, advisory support, API access, and licensing. It is an assumption-based operating case, not guaranteed revenue or audited guidance.

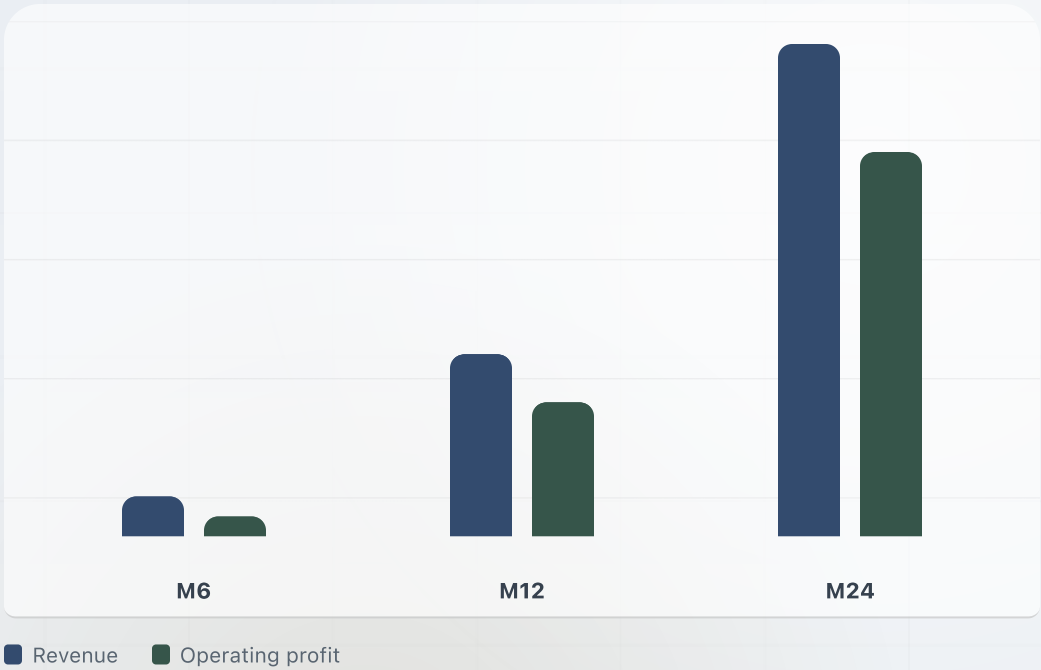

Checkpoints shown: M6 revenue $66K / profit $32.6K; M12 revenue $324.5K / profit $237.3K; M24 revenue $870K / profit $683K.

YEAR 1 REVENUE

$1.42M

Base scenario revenue across months 1-12.

YEAR 2 REVENUE

$7.22M

Base scenario revenue across months 13-24.

M24 RUN RATE

$10.44M

Month 24 revenue multiplied by 12.

BREAK-EVEN

M4

First modeled positive operating profit month.

09 / COMPUTE THESIS

The current bottleneck is compute.

The long-term moat is not simply the score interface. The current OS can evaluate smaller models; dedicated compute is required for larger-model benchmarking, reproducibility, and independent evaluation integrity.

Current OS machine

Already owned

Mac Studio M2 Max, 32GB. Enough to build the product, validate workflows, generate reports, and evaluate smaller local models.

- Workflow validation
- Console and report surfaces
- Prototype benchmark execution
- Constrained local inference

Founder lab compute

Approximately $17K-$20K

Mac Studio expansion plus two DGX Spark systems. Moves Norynthe into serious independent founder-scale evaluation.

- Larger open-weight models
- Repeatable benchmark execution
- Parallel evaluation runs
- Early customer-facing capacity

Production-grade unlock

Approximately $100K

GB300 workstation class infrastructure. Supports far larger open-weight workloads and higher-throughput evaluation cycles.

- Large-scale local inference
- Frontier-class open-weight workloads
- Enterprise-scale assessment pipelines
- Reduced API dependence

TECHNICAL CAVEAT

Capability ranges depend on model architecture, quantization, memory configuration, runtime optimization, context length, and evaluation workload complexity.

10 / DILIGENCE PATH

The next step is focused review.

Investors should be able to inspect the existing OS, understand the compute bottleneck, review the capital plan, and move into focused diligence with a clear sequence.

1

Open the OS preview

Review the console surfaces, record structure, report output, and smaller-model workflow architecture.

2

Review the raise and use of funds

Check whether the infrastructure-first capital plan matches the company thesis.

3

Read the financial assumptions

Evaluate the 24-month model as a base scenario and pressure-test the sales ramp.

4

Coordinate follow-up diligence

Schedule a focused review of the current OS, assumptions, capital plan, and open diligence questions.

Norynthe.

Diligence package.

OS preview, investor packet, financial model, raise plan, and supporting materials are organized for focused review.