

01 / EXECUTIVE SNAPSHOT

Norynthe is building the external trust layer for AI systems.

Norynthe.Score is the public signal. The Norynthe Model is the proprietary scoring engine that converts controlled benchmark assets into ranked model standing, score detail, and institution-ready reporting.

External evidence becomes the buying standard.

Enterprises need a way to compare AI systems outside model-owner claims, vendor demos, and narrow benchmark headlines. Norynthe creates a governed evidence record that can travel through procurement, risk, and executive decision processes.

1

Working MVP

Console, assessment records, public investor materials, and report surfaces are already in place.

2

Independent score layer

The score is built from governed benchmark assets, not model-owner marketing claims.

3

Trust lab unlock

The \$2M round funds controlled compute, evidence hardening, and enterprise validation.

01

Benchmark assets

Governed prompts, rubrics, test dimensions, and evaluation conditions.

02

Norynthe Model

Evaluator, scoring logic, confidence, blocker criteria, and judgment rules.

03

Assessment records

Versioned model standing, score detail, review state, findings, and flags.

04

Enterprise reports

Buyer-readable evidence for procurement, governance, risk, and operations.

02 / COMPANY THESIS

Independent scoring becomes infrastructure.

Most AI evaluation collapses into demos, brand perception, leaderboard headlines, or provider-controlled tests. Norynthe starts from a different premise: enterprises need a governed external standard.

CATEGORY THESIS

Trust moves outside the model.

Norynthe is building the record system enterprises can use when model claims are not enough.

Benchmark bank

Scoring engine

Assessment record

Report surface

01

AI choice is now an operating decision.

Model selection is moving into workflows with procurement, policy, legal, risk, and operational consequences.

02

Existing evaluation is too vendor-shaped.

Leaderboards and demos are useful, but they do not create a neutral record buyers can reuse across departments.

03

A governed evidence record becomes reusable infrastructure.

The score creates comparison pressure; the record and report create the paid evidence layer behind it.

04

The moat compounds through evidence density.

Benchmark governance, scoring logic, reviewer state, and buyer-facing reports harden together as more evaluations run.

03 / PRODUCT PROOF

The MVP is complete and working.

The investor preview shows the working console behind Norynthe.Score: governed evaluation launch, runtime trace, reviewer queue, stored assessment records, benchmark governance, and report-ready evidence.

Input, then create the assessment record.

01 **Evaluation surface**
Choose the governed benchmark set and scenario.

BENCHMARK SET

Core Governed Benchmark v2
BENCHMARK SCENARIO

Board memo on adopting an external

02 **Model target**
Confirm the evaluated model and runtime posture.

TARGET MODEL
Gemma 3 4B Instruct

Gemma 3 4B Instruct
Local / ollama
GEMMA 3 4B
RUNTIME ONLINE
Open model profile

03 **Output source**
Select how the target output enters the record.

Runtime generation
RUNTIME READY

Pasted output
OPERATOR PROVIDED

Imported transcript
EXTERNAL CAPTURE

Runtime generation: Ask the selected local runtime to generate the target output before scoring.

OPTIONAL TARGET OUTPUT
Optional. Leave blank to generate output from the selected runtime.

Model readiness, scoring queue, and record-writing path for the current setup.

Estimated run time
1m 23s
4B model + prompt + scoring dimensions. Actual runtime may shift with local load and response length.

Run timer
Starts after submit

Model
Gemma 3 4B Instruct

Runtime tag
gemma3:4b

Benchmark set
core_v2

Question
gov_001

■■■■ EVALUATION-RUN.PREVIEW

```
norynthe run --mode local-runtime \  
  --model "gemma3:4b" \  
  --benchmark "core_v2" \  
  --question "gov_001" \  
  --workbook "norynthe_v2"  
  
[wait] benchmark.load("core_v2")  
[wait] payload.question = "gov_001"  
[wait] runtime.ask("gemma3:4b",  
  payload.prompt)  
[wait] scorer.queue(primary_dimensions)  
[wait] critic.resolve(evidence, flags,  
  peer_outputs)  
[wait] records.write(assessment_record)
```

MVP PROOF

Working evaluation workflow, not a concept mockup.

Launch

Benchmark, target model, runtime posture, and evaluation conditions.

Score

Evaluator logic, dimensions, evidence checks, flags, and confidence.

Record

Assessment state, reviewer posture, persisted findings, and caveats.

Report

Buyer-facing output generated from the underlying evidence record.

NORYNTHE INVESTOR PACKET

04

04 / BUSINESS MODEL

Public trust scores create demand. Paid evidence access creates revenue.

Norynthe sells access to the evidence behind independent AI evaluation. The score is the entry point; revenue comes when teams need the methodology, comparisons, records, and reports behind it.

PUBLIC LAYER

Norynthe.Score

A clear public signal that creates comparison pressure and gives buyers a reason to ask for the evidence behind the ranking.

Full evaluation reports

Evidence trails

EVIDENCE LAYER

Reports and records

Paid access to full evaluation findings, methodology detail, score context, evidence trails, and decision-ready summaries.

Score methodology

Governance documentation

RECURRING LAYER

Enterprise access

Subscriptions, custom evaluations, vendor comparisons, advisory review, API access, and licensing as trust demands repeat.

Vendor comparisons

Custom evaluations

REVENUE LOGIC

The score creates the reason to ask. The paid product is the evidence behind the score: methodology, comparisons, documentation, and decision support.

05 / ENTERPRISE USE CASES

Norynthe turns model trust into a decision workflow.

The enterprise need is not simply more model access. It is better judgment about which systems should be used, under what constraints, with what evidence, and with what level of review.

COMPARE

Which model has standing?

Rank systems under a shared benchmark surface instead of vendor-specific claims.

EXPLAIN

Why did one system outrank another?

Expose dimensions, evidence, flags, and score context behind the public result.

ESCALATE

Where should human review remain mandatory?

Identify blocker criteria, caveats, risk posture, and sensitive-use constraints.

OPERATIONALIZE

How should the selected system be used?

Translate evaluation evidence into procurement, governance, and deployment language.

Model selection

Compare vendors under the same benchmark surface.

Procurement review

Reduce dependence on vendor-defined evidence.

Internal governance

Make trust judgments reviewable across teams.

Sensitive workflows

Apply blocker criteria and escalation logic.

Operational deployment

Decide where constraints and human oversight belong.

USE-CASE READ

The product is strongest where decisions need reviewable evidence, shared language, and a standard outside the model vendor.

06 / RAISE PLAN

\$2M to build the AI trust lab.

The round moves the company from convincing proof into a controlled evaluation lab and enterprise validation cycle that can survive customer scrutiny.

TARGET RAISE

\$2.0M

Infrastructure-first capital for an owned operating base, controlled compute, legal setup, security, and founder-led validation.

CAPITAL READ

The plan prioritizes durable lab capacity and validation before broad payroll expansion.

01

Secure the base

Property, legal setup, operating environment, and investor-ready governance foundation.

02

Build the lab

Workspace, controlled compute, storage, security, and repeatable evaluation infrastructure.

03

Harden the model

Benchmark governance, scoring logic, review-state discipline, and assessment records.

04

Validate with buyers

Enterprise workflows, pilots, paid report access, custom evaluation demand, and subscriptions.

07 / USE OF FUNDS

Capital goes into the lab before payroll scale.


The allocation prioritizes durable operating infrastructure, evaluation capacity, legal setup, controlled operations, business development, and founder runway.

TOTAL RAISE

\$2.0M

The use-of-funds view is intentionally literal: every bar is proportional to the full \$2M raise, so the dominant property and runway allocations remain visually honest.


Owned base	56.3%
Runway buffer	23.9%
Lab and compute	12.0%



08 / FINANCIAL MODEL

A 24-month base scenario after a six-month ramp.

The forecast models enterprise subscriptions, custom evaluations, reports, advisory support, API access, and licensing. It is an assumption-based operating case, not guaranteed revenue or audited guidance.



■ Revenue ■ Operating profit

Checkpoints shown: M6 revenue \$66K / profit \$32.6K; M12 revenue \$324.5K / profit \$237.3K; M24 revenue \$870K / profit \$683K.

YEAR 1 REVENUE

\$1.42M

Base scenario revenue across months 1-12.

YEAR 2 REVENUE

\$7.22M

Base scenario revenue across months 13-24.

M24 RUN RATE

\$10.44M

Month 24 revenue multiplied by 12.

BREAK-EVEN

M4

First modeled positive operating profit month.

09 / COMPUTE THESIS

Controlled evaluation requires controlled compute.

The long-term moat is not simply the score interface. It is controlled compute tied directly to benchmark governance, reproducibility, and independent evaluation integrity.

Current MVP machine

Already owned

Mac Studio M2 Max, 32GB. Enough to build the product, validate workflows, generate reports, and test smaller local models.

Workflow validation

Console and report surfaces

Prototype benchmark execution

Constrained local inference

Founder lab compute

Approximately \$17K-\$20K

Mac Studio expansion plus two DGX Spark systems. Moves Norynthe into serious independent founder-scale evaluation.

Larger open-weight models

Repeatable benchmark execution

Parallel evaluation runs

Early customer-facing capacity

Production-grade unlock

Approximately \$100K

GB300 workstation class infrastructure. Supports far larger open-weight workloads and higher-throughput evaluation cycles.

Large-scale local inference

Frontier-class open-weight workloads

Enterprise-scale assessment pipelines

Reduced API dependence

TECHNICAL CAVEAT

Capability ranges depend on model architecture, quantization, memory configuration, runtime optimization, context length, and evaluation workload complexity.

10 / DILIGENCE PATH

The next step is focused review.

Investors should be able to inspect the product, understand the capital plan, review the core assumptions, and move into focused diligence with a clear sequence.

1

Open the MVP preview

Review the console surfaces, record structure, report output, and workflow architecture.

2

Review the raise and use of funds

Check whether the infrastructure-first capital plan matches the company thesis.

3

Read the financial assumptions

Evaluate the 24-month model as a base scenario and pressure-test the sales ramp.

4

Coordinate follow-up diligence

Schedule a focused review of the MVP, assumptions, capital plan, and open diligence questions.

Norynthe.

Diligence package.

MVP preview, investor packet, financial model, raise plan, and supporting materials are organized for focused review.