

INDEPENDENT AI EVALUATION LABORATORY

Norynthe Investor Packet

Norynthe is building the controlled physical, computational, methodological, evidentiary, security, and governance environment required for independent AI evaluation.

RAISE TARGET

\$2.0M

Establish the controlled evaluation laboratory.

OPERATING MODEL

Founder-led

Low headcount by design; infrastructure and governed workflows carry the load.

CURRENT STATE

Foundation built

Published research, Method v0.1, architecture, prototype surfaces, and early experimentation.

80
SAMPLE SCORE

Build the institution that can earn the evidence.

The prototype makes the workflow concrete. The funded phase must produce complete run packages, repeatability results, external challenge, and traceable reports.

Output, then create the assessment record.

01 Evaluation surface
Choose the governed benchmark set and scenario.

BENCHMARK SET
Core Governed Benchmark v2

BENCHMARK SCENARIO
Board memo on adopting an external

02 Model target
Confirm the evaluated model and runtime posture.

TARGET MODEL
Gemma 3 4B Instruct

Gemma 3 4B Instruct
Local / ollama
GEMMA 3 4B
RUNTIME ONLINE
Open model profile

03 Output source
Select how the target output enters the record.

Runtime generation
RUNTIME READY

Pasted output
OPERATOR PROVIDED

Imported transcript
EXTERNAL CAPTURE

Runtime generation: Ask the selected local runtime to generate the target output before scoring.

OPTIONAL TARGET OUTPUT
Optional. Leave blank to generate output from the selected runtime.

Estimated run time
1m 23s

Run timer
Starts after submit

4B model + prompt + scoring dimensions. Actual runtime may shift with local load and response length.

Model
Gemma 3 4B Instruct

Runtime tag
gemma3:4b

Benchmark set
core_v2

Question
gov_001

EVALUATION-RUN.PREVIEW

```
norynthe run --mode local-runtime \
--model "gemma3:4b" \
--benchmark "core_v2" \
--question "gov_001" \
--workbook "norynthe_v2"

[wait] benchmark.load("core_v2")
[wait] payload.question = "gov_001"
[wait] runtime.ask("gemma3:4b",
payload.prompt)
[wait] scorer.queue(primary_dimensions)
[wait] critic.resolve(evidence, flags,
peer_outputs)
[wait] records.write(assessment_record)
```

OpenAI model family / enterprise decision support.

Established sub-model report showing the evidence behind a public Norynthe Score: score drivers, benchmark results, dimension behavior, review notes, and comparison context.

80

What the buyer

Capital Budget

The raise is modeled around durable infrastructure before payroll scale.

This model prioritizes fully owned operating infrastructure, evaluation capacity, legal advice, and nonpay. Instead of using the initial capital primarily for headcount, the budget is structured around the operating base required for independent AI evaluation.

USE OF FUNDS
\$2,000,000 capital deployment

PROPERTY / HQ infrastructure
\$115M
Lab / Workspaces Building
\$88.0K
Corporate Infrastructure
\$144.0K
Legal / Policy Setup
\$72K
Data / Software / Security
\$384K
Human Resources
\$500K
Compliance
\$477K
Remaining Liquidity / Buffer

REVENUE PROJECTIONS
Assumption table

CATEGORY	AMOUNT	%
Property / Headquarters Infrastructure	\$115,000,000	56.3%
Lab / Workspaces Building	\$88,000	4.3%
Corporate Infrastructure	\$144,000	7.2%
Legal / Policy Setup	\$72,000	3.6%
Data / Software / Security	\$384,000	1.9%
Human Resources	\$500,000	2.5%
Compliance	\$477,000	2.4%
Remaining Liquidity / Buffer	\$877,000	43.9%

01 / EXECUTIVE SNAPSHOT

Norynthe is building the independent laboratory for AI trust.

The laboratory joins controlled conditions, governed benchmarks and methods, hybrid compute, preserved evidence, external challenge, correction, and institutional reporting outside the model owner's control.

Independent evidence requires an institution built to produce it.

Enterprises need a way to compare AI systems outside model-owner claims, vendor demos, and narrow benchmark headlines. Norynthe creates a governed evidence record that can travel through procurement, risk, and executive decision processes.

1

Founder-built foundation

Published research, Method v0.1, prototype workflow surfaces, illustrative reports, and early experimentation make the design concrete.

2

Independent score layer

The score is built from governed benchmark assets, not model-owner marketing claims.

3

Laboratory capability unlock

The \$2M establishes the controlled base, hybrid compute, evidence systems, security, governance, external challenge, and founder-led runway.

01

Benchmark assets

Governed prompts, rubrics, test dimensions, and evaluation conditions.

02

Norynthe Model

Evaluator, scoring logic, confidence, blocker criteria, and judgment rules.

03

Assessment records

Versioned model standing, score detail, review state, findings, and flags.

04

Enterprise reports

Buyer-readable evidence for procurement, governance, risk, and operations.

02 / COMPANY THESIS

Independent scoring becomes infrastructure.

Most AI evaluation collapses into demos, brand perception, leaderboard headlines, or provider-controlled tests. Norynthe starts from a different premise: trustworthy evaluation requires an independent evidence-producing institution.

CATEGORY THESIS

Trust moves outside the model.

Norynthe is building the record system enterprises can use when model claims are not enough.

Benchmark bank

Scoring engine

Assessment record

Report surface

01

AI choice is now an operating decision.

Model selection is moving into workflows with procurement, policy, legal, risk, and operational consequences.

02

Existing evaluation is too vendor-shaped.

Leaderboards and demos are useful, but they do not create an independently governed evidence record buyers can reuse across departments.

03

A governed evidence record becomes reusable infrastructure.

The score creates comparison pressure; the record and report create the paid evidence layer behind it.

04

The moat compounds through evidence density.

Benchmark governance, scoring logic, reviewer state, and buyer-facing reports harden together as more evaluations run.

03 / CURRENT EVIDENCE STATE

A founder-built foundation; laboratory validation ahead.

The investor preview is a static prototype of the intended workflow. It is not a production backend or proof of completed scientific validation. The next proof is executable, preserved, repeatable run evidence.

...then create the assessment record.

01 **Evaluation surface**
Choose the governed benchmark set and scenario.

BENCHMARK SET

Core Governed Benchmark v2
BENCHMARK SCENARIO

Board memo on adopting an exten

02 **Model target**
Confirm the evaluated model and runtime posture.

TARGET MODEL
Gemma 3 4B Instruct

Gemma 3 4B Instruct
Local / ollama
GEMMA 3 4B
RUNTIME ONLINE
Open model profile

03 **Output source**
Select how the target output enters the record.

Runtime generation
RUNTIME READY

Pasted output
OPERATOR PROVIDED

Imported transcript
EXTERNAL CAPTURE

Runtime generation: Ask the selected local runtime to generate the target output before scoring.
OPTIONAL TARGET OUTPUT
Optional. Leave blank to generate output from the selected runtime.

Model inventory, scoring queue, and record-writing path for the current setup.

Estimated run time
1m 23s
4B model + prompt + scoring dimensions. Actual runtime may shift with local load and response length.

Run timer
Starts after submit

Model
Gemma 3 4B Instruct

Runtime tag
gemma3:4b

Benchmark set
core_v2

Question
gov_001

EVALUATION-RUN.PREVIEW

```
norynthe run --mode local-runtime \  
  --model "gemma3:4b" \  
  --benchmark "core_v2" \  
  --question "gov_001" \  
  --workbook "norynthe_v2"  
  
[wait] benchmark.load("core_v2")  
[wait] payload.question = "gov_001"  
[wait] runtime.ask("gemma3:4b",  
  payload.prompt)  
[wait] scorer.queue(primary_dimensions)  
[wait] critic.resolve(evidence, flags,  
  peer_outputs)  
[wait] records.write(assessment_record)
```

ILLUSTRATIVE PROTOTYPE

Intended evaluation workflow before controlled laboratory validation.

Launch

Benchmark, target model, runtime posture, and evaluation conditions.

Score

Evaluator logic, dimensions, evidence checks, flags, and confidence.

Record

Intended assessment state, reviewer posture, illustrative findings, and caveats.

Report

Buyer-facing output generated from the underlying evidence record.

NORYNTHE INVESTOR PACKET

04

04 / COMMERCIAL MODEL

A bounded paid workflow comes before recurring products.

The first commercial hypothesis is a scoped enterprise adoption-evidence engagement: one consequential decision, one controlled evaluation record, and one decision-ready report. Recurring products follow only after validation.

LABORATORY LAYER

Controlled evaluation

Governed conditions, run manifests, evidence preservation, review, limitations, and repeatability.

Full evaluation reports

Evidence trails

FIRST PAID WORKFLOW

Adoption-evidence engagement

A bounded evaluation record and report tied to one real institutional decision.

Score methodology

Governance documentation

LATER HYPOTHESES

Recurring evidence access

Subscriptions, monitoring, comparisons, APIs, or licensing only where repeat demand and delivery economics support them.

Vendor comparisons

Custom evaluations

REVENUE DISCIPLINE

Commercial forecasts are withheld until design partners establish willingness to pay, cycle length, delivery effort, decision impact, and repeat usage.

05 / ENTERPRISE USE CASES

Norynthe is designed to turn model trust into a decision workflow.

The funded laboratory will test whether controlled evidence improves judgment about which systems should be used, under what constraints, with what evidence, and with what level of review.

COMPARE

Which model has standing?

Test ranking under a shared benchmark surface instead of vendor-specific claims.

EXPLAIN

Why did one system outrank another?

Test whether dimensions, evidence, flags, and score context make a result inspectable.

ESCALATE

Where should human review remain mandatory?

Test blocker criteria, caveats, risk posture, and sensitive-use constraints.

OPERATIONALIZE

How should the selected system be used?

Test whether evidence supports procurement, governance, and deployment decisions.

Model selection

Compare vendors under the same benchmark surface.

Procurement review

Reduce dependence on vendor-defined evidence.

Internal governance

Make trust judgments reviewable across teams.

Sensitive workflows

Apply blocker criteria and escalation logic.

Operational deployment

Decide where constraints and human oversight belong.

USE-CASE HYPOTHESIS

Norynthe should be strongest where decisions need reviewable evidence, shared language, and an evaluator outside the model vendor's control.

06 / RAISE PLAN

\$2M to establish the controlled evaluation laboratory.

The round creates the integrated physical, computational, methodological, evidentiary, security, governance, and operating capacity required for repeatable independent evaluation.

TARGET RAISE

\$2.0M

Laboratory-first capital for a controlled base, hybrid compute, evidence and security systems, method governance, external challenge, and founder-led runway.

OPERATING READ

The plan remains intentionally low-headcount. Capital builds reusable capability and institutional evidence before payroll expansion.

01

Specify and govern

Facility alternatives, compute configuration, method register, benchmark register, security, conflicts, and correction.

02

Commission controlled execution

Minimum viable environment, evidence ledger, version controls, chain of custody, and first complete run package.

03

Demonstrate repeatability

Repeated-run variance, method reconciliation, external challenge, and traceable pilot reports.

04

Validate institutional demand

Bounded design-partner engagements, willingness to pay, decision impact, and supported recurring paths.

07 / USE OF FUNDS

Capital goes into governed laboratory capability before payroll scale.

The allocation establishes a controlled base, hybrid compute, evidence and security systems, governance, institutional validation, and founder-led runway. The largest category remains subject to explicit facility and ownership diligence.

TOTAL RAISE

\$2.0M

The use-of-funds view is intentionally literal. The controlled-base assumption is dominant and therefore must pass functional necessity, technical suitability, economic, governance, and capital-efficiency tests before commitment.


Controlled base	56.3%
Founder-led runway	23.9%
Lab and compute	12.0%

**Controlled laboratory base —
diligence gated**
\$1,126,571




56.3%

**Founder-led laboratory
runway**
\$477,429




23.9%

Compute infrastructure
\$146,500




7.3%

Lab / workspace buildout
\$93,500




4.7%

**Legal / entity / IP / property
governance**
\$72,000




3.6%

**Institutional validation /
design partners**
\$50,000



2.5%

Ops / software / security
\$34,000



1.7%

08 / 18-MONTH PROOF PROGRAM

Capital deployment is staged against evidence, not calendar optimism.

Prior internal revenue scenarios are not presented as investor guidance. The funded phase retires technical, methodological, institutional, and commercial risk in sequence.

00–03 Govern the foundation

One canonical method and benchmark relationship; facility, entity, IP, compute, evidence, security, conflict, appeal, and correction diligence.

03–06 Commission the minimum viable laboratory

Controlled execution environment, evidence ledger, version controls, chain of custody, and first complete verified run package.

06–12 Earn repeatability

Repeated-run variance results, reconciled method and benchmark versions, external methodological challenge, and traceable pilot reports.

12–18 Validate institutional demand

Bounded design-partner engagements, willingness to pay, decision impact, delivery economics, and supported recurring products.

FORECAST GATE

A board-usable commercial forecast follows evidence on cycle length, pricing, gross delivery effort, repeat usage, conversion, and the effect of Norynthe's evidence on a real decision.

09 / LABORATORY ARCHITECTURE

Compute matters inside a governed evidence system.

No workstation creates independence by itself. Norynthe requires integrated control across conditions, methods, benchmark custody, execution, evidence, security, review, correction, and reporting.

Method and evidence controls

Institutional core

Canonical methods, governed benchmarks, run manifests, evidence ledger, chain of custody, review, correction, and institutional memory.

Traceability

Version control

Challenge and appeal

Preserved limitations

Hybrid compute

Configuration under diligence

Selected local capacity, secure storage, APIs, and external capacity chosen by control, coverage, repeatability, security, and cost.

Local sensitive work

Frontier API coverage

Repeated execution

Cost discipline

Controlled physical base

Diligence gated

A purpose-bound environment for evidence custody, security, infrastructure continuity, protected benchmark work, and founder-led operation.

Power and thermal control

Physical security

Operating continuity

Company governance

FACILITY CAVEAT

The dominant base allocation must pass functional necessity, technical suitability, buy-versus-lease economics, company-control governance, and capital-efficiency tests before commitment.

10 / FOUNDER AND DILIGENCE PATH

A low-headcount institution with explicit founder accountability.

Alan Motley remains the principal operator, laboratory director, method steward, and institutional lead. His standpoint as a Black AI systems strategist informs which questions Norynthe asks; governed evidence must earn every finding.

1

Inspect the evidence state

Separate published method, prototype surfaces, synthetic samples, early experiments, and capital-enabled proof.

2

Review laboratory necessity

Test whether each physical, computational, evidentiary, security, and governance component produces a capability the institution needs.

3

Challenge the capital map

Pressure-test facility alternatives, hardware configuration, governance, runway, milestone gates, and the appropriate blend of capital.

4

Define the next proof

Agree on complete run packages, repeatability, external challenge, design partners, paid pilots, and decision impact.

Norynthe.

Laboratory-first diligence.

The packet is organized to answer why the laboratory must exist, why it remains founder-led and low-headcount, what capital establishes, and what evidence the funded phase must earn.

investors.norynthe.com

hello@norynthe.com